

Advanced Databases

3

Reliability Control

Topics


- Persistence of memory and backup
- Buffer management
- Reliable transaction management
- Log management
- Recovery after malfunctions

Persistence of Memories

- Main memory
 - Not persistent
- Mass memory
 - Persistent but can be damaged
- Stable memory
 - Cannot be damaged (it's an abstraction)


How to Guarantee Stable Memory

- On-line replication: mirroring of two disks




How to Guarantee Stable Memory

- Off-line replication: tape unit (backup)




Main Memory Management



- Rationale
 - Reuse of data in the buffer
 - Deferred writing into the database

Use of Main Memory



Buffer Management

- Based on four primitives:
 - **fix**
 - Used to load a page into the buffer. After the operation, the page is allocate to a transaction
 - **use**
 - Use of a page in the buffer
 - **unfix**
 - De-allocation of a page
 - **force**
 - Synchronously transfers a page from the main memory to the database

Using the Buffer (Transactions)

- Follows the scheme
fix
repeat **use** until (end of transaction)
unfix
- Pages are written by the buffer asynchronously
- **flush**
 - This primitive is controlled by the buffer manager
 - Asynchronously transfers a page from the main memory to the database

Executing a *fix* Primitive

- Searching for the target page
 - Selection of a free page
 - Otherwise, selection of a de-allocated page, which, if necessary, is copied onto the disk
 - Otherwise (if **STEAL** policy) a page is taken away from an active transaction. The page is copied onto the disk
 - Otherwise (if **NO STEAL** policy) the search fails
- Reading
 - If a target page exists, it is read from the database into the buffer in main memory

Buffer Management Policies


- STEAL (pages taken away from an active transaction)
- NO STEAL
- FORCE (pages written at commit-work)
- NO FORCE
- Normally:
 - NO STEAL
 - NO FORCE

Buffer Management Policies

- PRE-FETCHING
 - anticipates reading of pages
 - especially useful in sequential reading
- PRE-FLUSHING
 - anticipates writing of de-allocated pages
 - useful for accelerating page **fix**


Reminder: Atomicity Requirements

- A transaction is an atomic transformation from an initial state into a final state
- Possible behaviors:
 - `commit work`: success
 - `rollback work` or error before `commit`: undo
 - Fault after `commit`: redo



Transaction Log

- Sequential file consisting of records that describe the actions carried out by the various transactions
- Written sequentially to the top block (top = current instant)



Main Function of the Log

- It records in the stable memory the actions carried out by the various transactions under the form of state transitions

If **UPDATE (U)**

transforms **O** from value **O1** to value **O2**

then the log records:

BEFORE-STATE (U) = O1

AFTER-STATE (U) = O2

Using the Log

- After `rollback-work` or failure
 - `UNDO T1: 0 = 01`
- After failure after `commit`
 - `REDO T1: 0 = 02`
- Idempotency of `UNDO` and `REDO`:
$$\text{UNDO}(T) = \text{UNDO}(\text{UNDO}(T))$$
$$\text{REDO}(T) = \text{REDO}(\text{REDO}(T))$$

Types of Log Records

- Records relevant to transactional commands:
 - `begin`
 - `commit`
 - `abort`
- Records relevant to operations
 - `insert`
 - `delete`
 - `update`
- Records relevant to recovery actions
 - `dump`
 - `checkpoint`


Types of Log Records

- Records relevant to transactional commands:
 - $B(T), C(T), A(T)$
- Records relevant to operations
 - $I(T, O, AS), D(T, O, BS), U(T, O, BS, AS)$
- Records relevant to recovery actions
 - $DUMP, CKPT(T_1, T_2, \dots, T_n)$
- Record fields:
 - T : transaction identifier
 - O : object identifier
 - BS, AS : before state, after state

Transactional Rules


- Write-Ahead-Log
 - Before-state parts of the log records must be written in the log before carrying out the corresponding operation on the database
 - Actions can be undone
- Commit Rule
 - After-state parts of the log records must be written in the log before carrying out the commit
 - Actions can be redone

Writing onto Log and Database



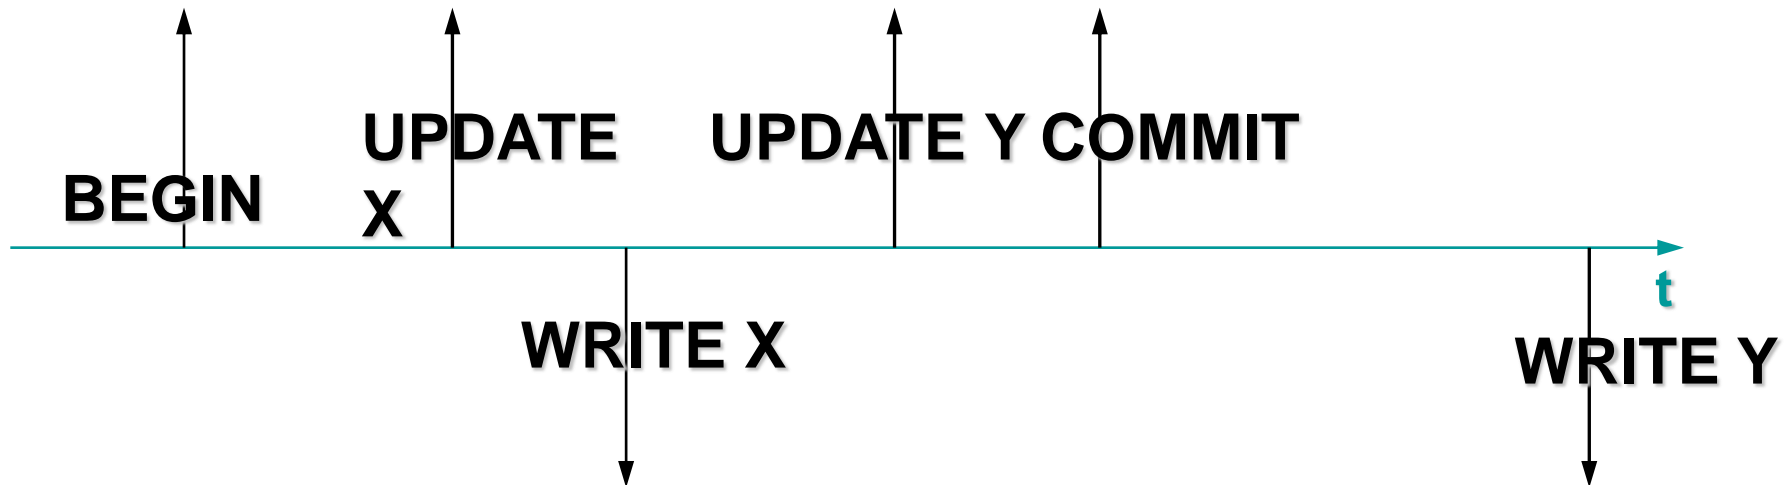
- Writing onto the database before commit
 - Requires writing in order to abort

Writing onto Log and Database



- Writing onto the database after commit
 - Does not require writing in order to abort

Writing onto Log and Database

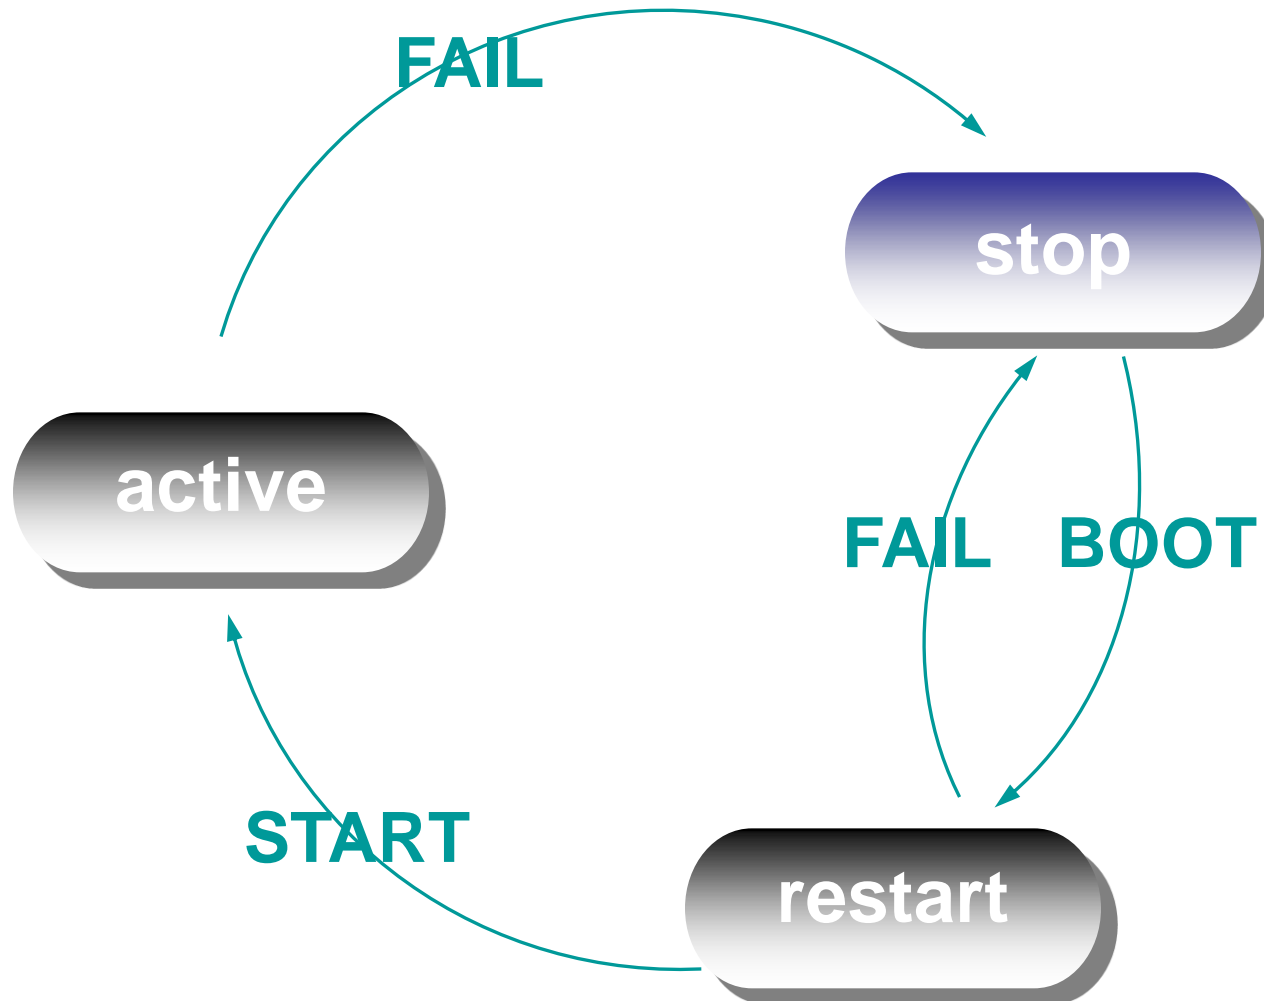


- Writing onto the database in an arbitrary moment
 - Allows optimizing buffer management

In Case of Failure


- Soft failure
 - Loss of the contents of the main memory
 - Requires warm restart
- Hard failure
 - Failure of secondary memory devices
 - Requires cold restart

Fail-stop Failure Model



Checkpoint

- “Consistent” time point
(in which all transactions write their data from the buffer to the disk)
- All active transactions are recorded



Checkpoint

- Operation used to “sum things up”, by simplifying the subsequent restore operations
 - Aim: to record which transactions are active at a given moment (and, dually, to confirm that the others either did not start or have finished)
- Parallel (extreme):
 - Closing the balance at the end of the year
 - Example: since November 25 no new “operation” request is accepted and all previously initiated operations must be concluded before new ones can be accepted


Checkpoint

- Several possibilities – the simplest is as follows:
 1. Acceptance of commit requests is suspended.
 2. All dirty pages written by committed transactions are transferred to mass storage (via **force**).
 3. The identifiers of the transactions in progress are recorded on the log (via **force**); no new transaction can start while this recording takes place.

Then, acceptance of operations is resumed
- This way, we are sure that
 - For all committed transactions, the data are on mass storage
 - Transactions that are “half-way” are listed in the checkpoint

Dump


- Time point in which a complete copy of the database is created (typically during the night or the week-end)
- The presence of the dump is recorded



Warm Restart

- Log records are read starting from the checkpoint
- Transactions are divided into:
 - UNDO set
 - REDO set
- UNDO and REDO actions are executed

Warm restart



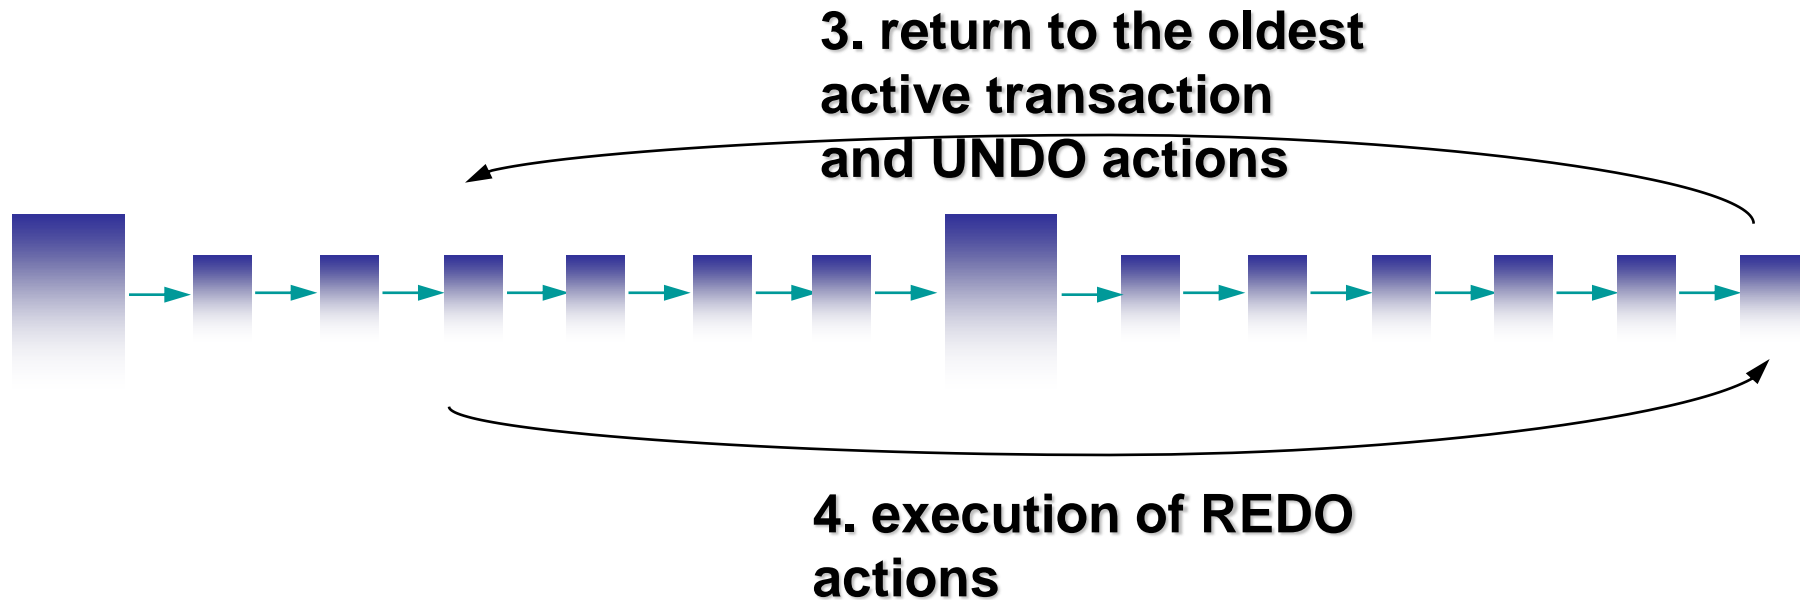
- UNDO

- Active transactions before commit

- REDO

- Active transactions after commit

Warm Restart



Example of Warm Restart

- B(T1)
 - B(T2)
 - U(T1,O1,B1,A1)
 - I(T1,O2,A2)
 - U(T2,O3,B3,A3)
 - B(T3)
 - U(T3,O4,B4,A4)
 - D(T3,O5,B5)
 - CKPT(T1,T2,T3)
 - C(T2)
 - B(T4)
 - U(T4,O6,B6,A6)
 - A(T4)
 - failure
- UNDO=(T1,T2,T3)
REDO=()
 - UNDO=(T1,T3,T4)
REDO=(T2)

Example of Warm Restart

- B(T1)
 - B(T2)
 - U(T1,O1,B1,A1)
 - I(T1,O2,A2)
 - U(T2,O3,B3,A3)
 - B(T3)
 - U(T3,O4,B4,A4)
 - D(T3,O5,B5)
 - CKPT(T1,T2,T3)
 - C(T2)
 - B(T4)
 - U(T4,O6,B6,A6)
 - A(T4)
 - failure
- UNDO=(T1,T3,T4)
REDO=(T2)
- O1 = B1
 - DELETE(O2)
 - O3 = A3
 - O4 = B4
 - O5 = B5
 - O6 = B6
- RESTART

Cold Restart

- Data are restored starting from the backup
- The operations recorded onto the log until the failure are executed
- A warm restart is executed

Architecture of the Reliability Manager

